Attention Is All You Need

Ashish Vaswani, Noam Shazeer, Niki Parmar

Year 2017 doi:10.48550/arXiv.1706.03762

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely.

1. Introduction

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

2. Method

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

References

[1] J. Devlin, M. Chang. BERT: Pre-training of Deep Bidirectional Transformers for Language Und erstanding. NAACL 2019.

[2] A. Radford. Improving Language Understanding by Generative Pre-Training. OpenAI 2018.

[3] I. Sutskever. Sequence to Sequence Learning with Neural Networks. NeurIPS 2014.