BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin, Ming-Wei Chang, Kenton Lee

Year 2019 doi:10.18653/v1/N19-1423

Abstract

We introduce a new language representation model called BERT, which stands for Bidirectional Encoder Representations from Transformers. BERT obtains new state-of-the-art results on eleven natural language processing tasks.

1. Introduction

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

2. Method

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

References

[1] A. Vaswani, N. Shazeer. Attention Is All You Need. NeurIPS 2017. doi:10.48550/arXiv.1706.03762

[2] A. Radford. Improving Language Understanding by Generative Pre-Training. OpenAI 2018.