Deep Residual Learning for Image Recognition

Kaiming He, Xiangyu Zhang, Shaoqing Ren

Year 2016 doi:10.1109/CVPR.2016.90

Abstract

Deeper neural networks are more difficult to train. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We explicitly reformulate the layers as learning residual functions.

1. Introduction

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

2. Method

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

References