Sequence to Sequence Learning with Neural Networks

Ilya Sutskever, Oriol Vinyals, Quoc V. Le

Year 2014 doi:10.48550/arXiv.1409.3215

Abstract

This paper presents a general end-to-end approach to sequence learning that makes minimal assumptions on the sequence structure. Our method uses a multilayered Long Short-Term Memory (LSTM) to map the input sequence to a vector of a fixed dimensionality.

1. Introduction

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

2. Method

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

References