An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale

Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov

Year 2021 doi:10.48550/arXiv.2010.11929

Abstract

While the Transformer architecture has become the de-facto standard for natural language processing tasks, its applications to computer vision remain limited. We show that a pure transformer applied directly to sequences of image patches can perform very well on image classification tasks.

1. Introduction

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

2. Method

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

This paper studies the problem in a principled way. We describe the method, the theory behind it, and extensive experiments that validate the approach. The results show clear improvements over prior work across benchmarks.

References

[1] A. Vaswani, N. Shazeer, N. Parmar. Attention Is All You Need. NeurIPS 2017.

[2] K. He, X. Zhang, S. Ren. Deep Residual Learning for Image Recognition. CVPR 2016. doi:10.1109/CVPR.2016.90

[3] J. Devlin. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. 2019.